

# Predicting Housing Prices in Delhi, India Using Neural Networks

House Price Dataset of Delhi, scraped from Magic Bricks between June 2021 and July 2021, containing 18 columns.

```
In [1]: !kaggle datasets download -d goelyash/housing-price-dataset-of-delhiindia
Dataset URL: https://www.kaggle.com/datasets/goelyash/housing-price-dataset-of-delhiindia
License(s): CC0-1.0
housing-price-dataset-of-delhiindia.zip: Skipping, found more recently modified local copy (use --force to force download)

In [2]: # Surprise warnings:
def warn(*args, **kwargs):
    pass
import warnings
warnings.warn = warn

In [3]: import zipfile
zip_ref=zipfile.ZipFile("housing-price-dataset-of-delhiindia.zip", "r")
zip_ref.extractall()
zip_ref.close()

In [4]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

In [5]: df=pd.read_csv("Delhi_v2.csv")
df.head()

Out[5]:
```

	Unnamed: 0	price	Address	area	latitude	longitude	Bedrooms	Bathrooms	Balcony	Status	neworold	parking	Furnished_stat
0	0	5600000.0	Noida Extension, Noida, Delhi NCR	1350.0	28.608850	77.460560	3.0	3.0	NaN	Under Construction	New Property	NaN	N
1	1	8800000.0	Sector 79, Gurgaon, Delhi NCR	1490.0	28.374236	76.952416	3.0	3.0	NaN	Ready to Move	New Property	NaN	Semi-Furnish
2	2	16500000.0	Vaishali, Ghaziabad, Delhi NCR	2385.0	28.645769	77.385110	4.0	5.0	NaN	Ready to Move	New Property	1.0	Unfurnish
3	3	3810000.0	Link Road, F Block, Sector 50, Noida, Uttar Pr...	1050.0	28.566914	77.436434	2.0	2.0	3.0	NaN	New Property	1.0	Unfurnish
4	4	6200000.0	Jaypee Pavilion Court Sector 128, Noida, Secto...	1350.0	28.520732	77.356491	2.0	2.0	3.0	Ready to Move	Resale	1.0	N

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7738 entries, 0 to 7737
Data columns (total 18 columns):
 #   Column            Non-Null Count  Dtype  
 ---  -- 
 0   Unnamed: 0         7738 non-null   int64  
 1   price              7738 non-null   float64 
 2   Address             7738 non-null   object  
 3   area                7738 non-null   float64 
 4   latitude             7738 non-null   float64 
 5   longitude            7738 non-null   float64 
 6   Bedrooms             7738 non-null   float64 
 7   Bathrooms            7738 non-null   float64 
 8   Balcony              5166 non-null   float64 
 9   Status               7164 non-null   object  
 10  neworold             7738 non-null   object  
 11  parking              2612 non-null   float64 
 12  Furnished_status     4124 non-null   object  
 13  Lift                 1733 non-null   float64 
 14  Landmarks             2759 non-null   object  
 15  type_of_building      7738 non-null   object  
 16  desc                  7738 non-null   object  
 17  Price_sqft            7738 non-null   float64 
dtypes: float64(10), int64(1), object(7)
memory usage: 1.1+ MB
```

In [7]: df.describe()

	Unnamed: 0	price	area	latitude	longitude	Bedrooms	Bathrooms	Balcony	parking	Lift	Price
<b>count</b>	7738.000000	7.738000e+03	7738.000000	7738.000000	7738.000000	7738.000000	7738.000000	5166.000000	2612.000000	1733.000000	7738.00
<b>mean</b>	3868.500000	8.320635e+06	1409.506591	28.552092	77.273476	2.708193	2.501163	2.426442	7.673047	1.829198	5543.66
<b>std</b>	2233.912524	7.223197e+06	718.929581	0.107420	0.180606	0.877026	0.867050	1.083677	60.417508	1.192607	2408.65
<b>min</b>	0.000000	1.700000e+06	501.000000	28.240023	76.884101	2.000000	2.000000	1.000000	1.000000	0.000000	2100.00
<b>25%</b>	1934.250000	4.200000e+06	990.000000	28.455539	77.078590	2.000000	2.000000	2.000000	1.000000	1.000000	3950.98
<b>50%</b>	3868.500000	6.000000e+06	1250.000000	28.574637	77.345320	3.000000	2.000000	2.000000	1.000000	2.000000	4972.67
<b>75%</b>	5802.750000	9.500000e+06	1650.000000	28.642520	77.421054	3.000000	3.000000	3.000000	2.000000	2.000000	6350.63
<b>max</b>	7737.000000	8.500000e+07	9500.000000	28.799748	77.688028	10.000000	10.000000	10.000000	1406.000000	10.000000	44378.69

In [8]: df.nunique()

```
Out[8]: Unnamed: 0          7738
price              584
Address             4145
area                1102
latitude            2942
longitude           2912
Bedrooms             9
Bathrooms            9
Balcony              10
Status                2
neworold              2
parking              65
Furnished_status      3
Lift                 11
Landmarks             2612
type_of_building       2
desc                  4181
Price_sqft            3946
dtype: int64
```

In [9]: df.corr(numeric\_only=True)

Out[9]:

	Unnamed: 0	price	area	latitude	longitude	Bedrooms	Bathrooms	Balcony	parking	Lift	Price_sqft
<b>Unnamed: 0</b>	1.000000	0.022589	0.002100	-0.030258	0.008796	-0.004658	0.008812	0.047535	-0.019743	0.006369	0.036812
<b>price</b>	0.022589	1.000000	0.849073	-0.227517	-0.331883	0.606914	0.691990	0.255096	-0.041640	0.037892	0.682768
<b>area</b>	0.002100	0.849073	1.000000	-0.279069	-0.184427	0.683479	0.763650	0.420476	-0.036780	0.128409	0.292711
<b>latitude</b>	-0.030258	-0.227517	-0.279069	1.000000	0.401701	-0.140817	-0.246663	-0.199374	-0.053874	-0.094808	-0.098981
<b>longitude</b>	0.008796	-0.331883	-0.184427	0.401701	1.000000	-0.166826	-0.205187	0.090932	0.019576	0.139778	-0.451146
<b>Bedrooms</b>	-0.004658	0.606914	0.683479	-0.140817	-0.166826	1.000000	0.765479	0.283880	-0.037920	-0.048828	0.348992
<b>Bathrooms</b>	0.008812	0.691990	0.763650	-0.246663	-0.205187	0.765479	1.000000	0.361291	-0.043036	0.068791	0.356413
<b>Balcony</b>	0.047535	0.255096	0.420476	-0.199374	0.090932	0.283880	0.361291	1.000000	0.000870	0.331883	-0.012863
<b>parking</b>	-0.019743	-0.041640	-0.036780	-0.053874	0.019576	-0.037920	-0.043036	0.000870	1.000000	0.037927	-0.045357
<b>Lift</b>	0.006369	0.037892	0.128409	-0.094808	0.139778	-0.048828	0.068791	0.331883	0.037927	1.000000	-0.133483
<b>Price_sqft</b>	0.036812	0.682768	0.292711	-0.098981	-0.451146	0.348992	0.356413	-0.012863	-0.045357	-0.133483	1.000000

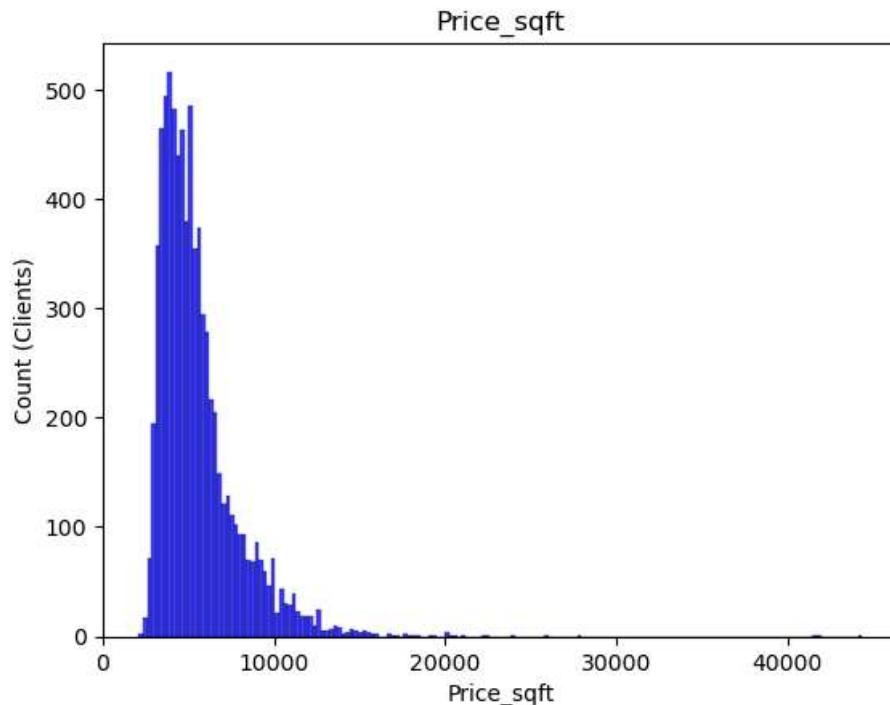
In [10]:

```
df = df.drop(["Address", "Landmarks", "price", "Unnamed: 0", "desc"], axis=1)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7738 entries, 0 to 7737
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   area             7738 non-null   float64 
 1   latitude         7738 non-null   float64 
 2   longitude        7738 non-null   float64 
 3   Bedrooms         7738 non-null   float64 
 4   Bathrooms        7738 non-null   float64 
 5   Balcony          5166 non-null   float64 
 6   Status            7164 non-null   object  
 7   neworold          7738 non-null   object  
 8   parking           2612 non-null   float64 
 9   Furnished_status 4124 non-null   object  
 10  Lift              1733 non-null   float64 
 11  type_of_building 7738 non-null   object  
 12  Price_sqft       7738 non-null   float64 
dtypes: float64(9), object(4)
memory usage: 786.0+ KB
```

In [11]:

```
sns.histplot(data=df, x="Price_sqft", color="b")
plt.xlabel("Price_sqft")
plt.ylabel("Count (Clients)")
plt.title("Price_sqft")
plt.show()
pd.DataFrame(df["Price_sqft"].describe())
```



```
Out[11]:
```

	Price_sqft
count	7738.000000
mean	5543.660241
std	2408.659307
min	2100.000000
25%	3950.986915
50%	4972.674332
75%	6350.638693
max	44378.698225

```
In [12]: out=dff["Price_sqft"].quantile(0.95)  
out
```

```
Out[12]: 10000.0
```

Removing the outliers

```
In [13]: datf=dff[dff["Price_sqft"]<=out]  
datf.reset_index(inplace=True,drop=True)  
datf.shape
```

```
Out[13]: (7360, 13)
```

```
In [14]: datf.head()
```

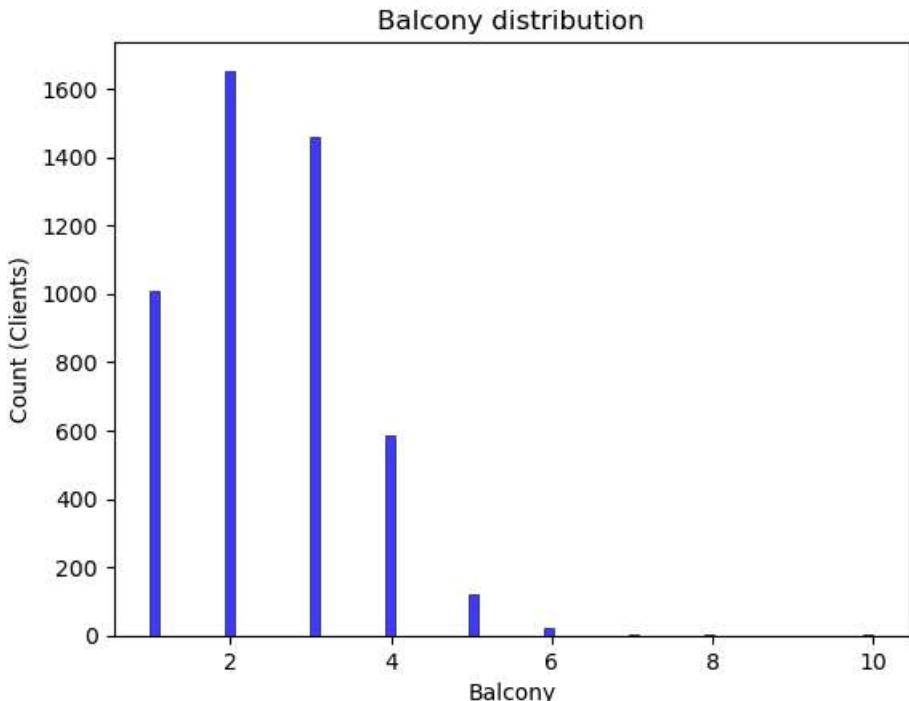
Out[14]:

	area	latitude	longitude	Bedrooms	Bathrooms	Balcony	Status	neworold	parking	Furnished_status	Lift	type_of_building	Price_
0	1350.0	28.608850	77.460560	3.0	3.0	NaN	Under Construction	New Property	NaN	NaN	2.0	Flat	4148.148
1	1490.0	28.374236	76.952416	3.0	3.0	NaN	Ready to Move	New Property	NaN	Semi-Furnished	2.0	Flat	5906.040
2	2385.0	28.645769	77.385110	4.0	5.0	NaN	Ready to Move	New Property	1.0	Unfurnished	NaN	Flat	6918.238
3	1050.0	28.566914	77.436434	2.0	2.0	3.0	NaN	New Property	1.0	Unfurnished	2.0	Flat	3628.571
4	1350.0	28.520732	77.356491	2.0	2.0	3.0	Ready to Move	Resale	1.0	NaN	3.0	Flat	4592.592

Replacing missing values

In [15]:

```
sns.histplot(data=datf,x="Balcony",color="b")
plt.xlabel("Balcony")
plt.ylabel("Count (Clients)")
plt.title("Balcony distribution")
plt.show()
print(pd.DataFrame(datf["Balcony"].value_counts()))
```



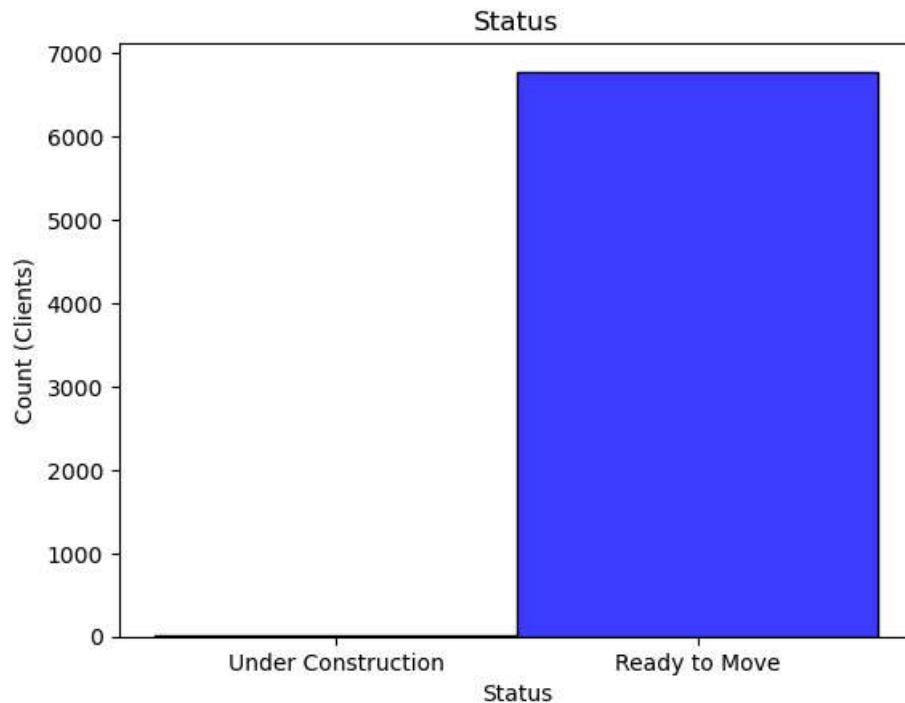
Balcony	count
2.0	1653
3.0	1457
1.0	1007
4.0	588
5.0	121
6.0	20
10.0	3
7.0	3
8.0	3
9.0	1

In [16]:

```
datf["Balcony"].fillna(df["Balcony"].mode().iloc[0],inplace=True)
```

In [17]:

```
sns.histplot(data=datf,x="Status",color="b")
plt.xlabel("Status")
plt.ylabel("Count (Clients)")
plt.title("Status")
plt.show()
print(pd.DataFrame(datf["Status"].value_counts()))
```



```
count
Status
Ready to Move      6782
Under Construction     14
```

```
In [18]: datf["Status"].fillna(df["Status"].mode().iloc[0], inplace=True)
```

```
In [19]: pd.DataFrame(datf["parking"].value_counts())
```

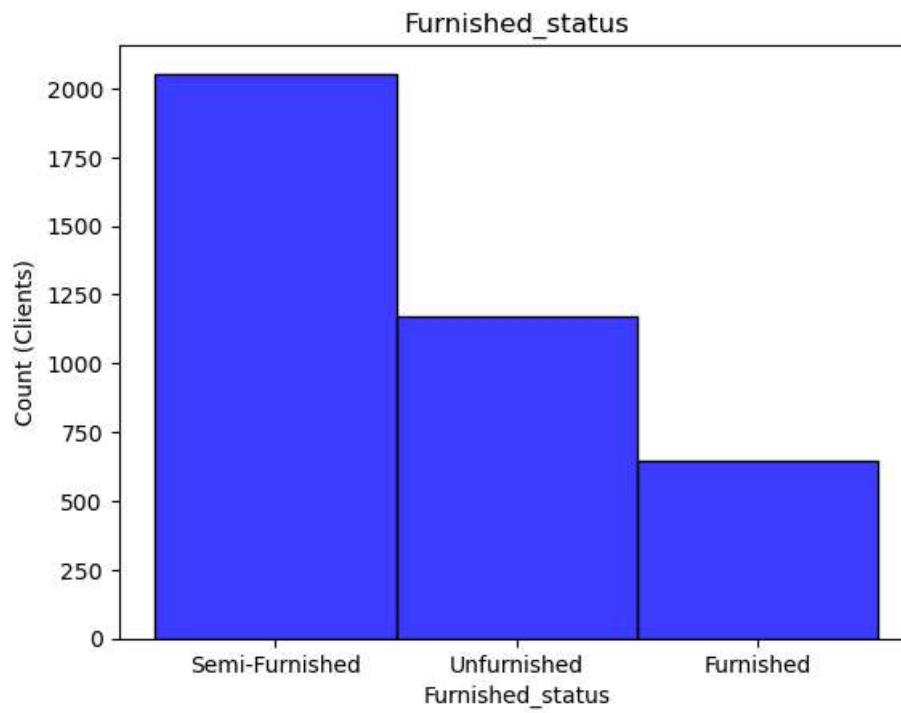
```
Out[19]:    count
```

parking	count
1.0	1651
2.0	574
3.0	84
4.0	29
6.0	10
...	...
18.0	1
602.0	1
16.0	1
624.0	1
179.0	1

63 rows × 1 columns

```
In [20]: datf["parking"].fillna(df["parking"].mode().iloc[0], inplace=True)
```

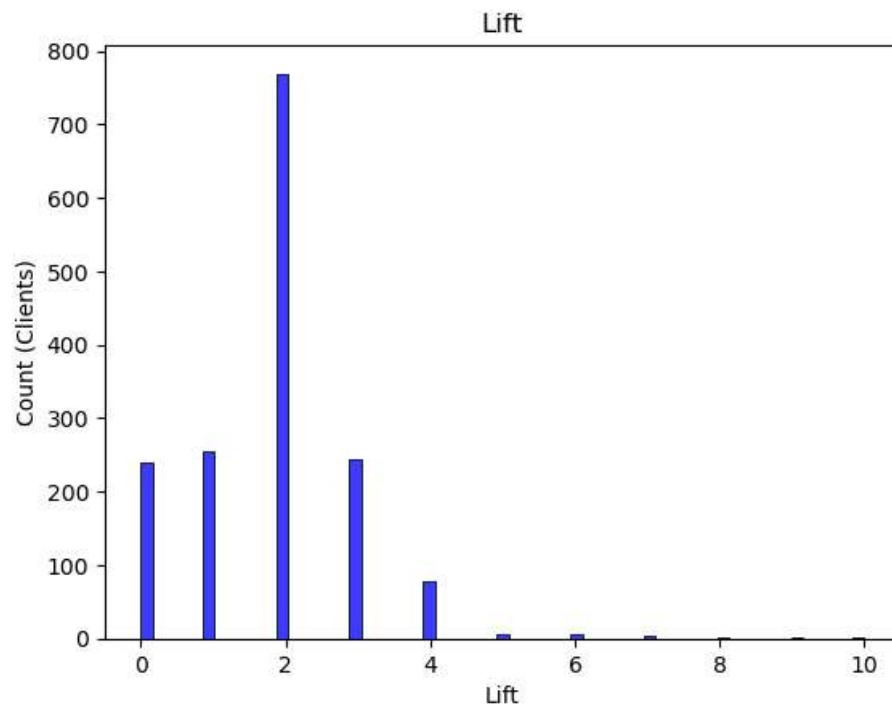
```
In [21]: sns.histplot(data=datf, x="Furnished_status", color="b")
plt.xlabel("Furnished_status")
plt.ylabel("Count (Clients)")
plt.title("Furnished_status")
plt.show()
print(pd.DataFrame(datf["Furnished_status"].value_counts()))
```



```
count
Furnished_status
Semi-Furnished      2055
Unfurnished         1172
Furnished           644
```

```
In [22]: datf["Furnished_status"].fillna(df["Furnished_status"].mode().iloc[0], inplace=True)
```

```
In [23]: sns.histplot(data=datf,x="Lift",color="b")
plt.xlabel("Lift")
plt.ylabel("Count (Clients)")
plt.title("Lift")
plt.show()
print(pd.DataFrame(datf["Lift"].value_counts()))
```



```
count
Lift
2.0    769
1.0    254
3.0    245
0.0    240
4.0     77
6.0      7
5.0      6
7.0      4
10.0     2
8.0      1
9.0      1
```

```
In [24]: datf["Lift"].fillna(df["Lift"].mode().iloc[0],inplace=True)
```

```
In [25]: datf.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7360 entries, 0 to 7359
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   area             7360 non-null   float64
 1   latitude         7360 non-null   float64
 2   longitude        7360 non-null   float64
 3   Bedrooms         7360 non-null   float64
 4   Bathrooms        7360 non-null   float64
 5   Balcony          7360 non-null   float64
 6   Status            7360 non-null   object  
 7   neworold         7360 non-null   object  
 8   parking           7360 non-null   float64
 9   Furnished_status 7360 non-null   object  
 10  Lift              7360 non-null   float64
 11  type_of_building 7360 non-null   object  
 12  Price_sqft       7360 non-null   float64
dtypes: float64(9), object(4)
memory usage: 747.6+ KB
```

```
In [26]: datf.nunique()
```

```
area            1057
latitude        2719
longitude       2694
Bedrooms         9
Bathrooms        9
Balcony          10
Status            2
neworold          2
parking           63
Furnished_status 3
Lift              11
type_of_building 2
Price_sqft        3715
dtype: int64
```

Converting the categorical entries into numerical ones

```
from sklearn.preprocessing import LabelEncoder
final_df=datf
for column in final_df.columns:
    if final_df[column].dtype == 'object':
        le = LabelEncoder()
        final_df[column] = le.fit_transform(final_df[column])
```

```
In [28]: final_df.corr(numeric_only=True)
```

Out[28]:

	area	latitude	longitude	Bedrooms	Bathrooms	Balcony	Status	neworold	parking	Furnished_status	Lift	typ
<b>area</b>	1.000000	-0.276648	-0.171144	0.708508	0.773638	0.361521	0.024622	-0.067268	-0.013302	-0.014781	0.043350	
<b>latitude</b>	-0.276648	1.000000	0.422607	-0.144976	-0.248055	-0.151663	-0.005987	0.057831	-0.036681	-0.044814	-0.023581	
<b>longitude</b>	-0.171144	0.422607	1.000000	-0.156850	-0.194897	0.059638	0.004647	0.097625	0.003697	-0.008245	0.077498	
<b>Bedrooms</b>	0.708508	-0.144976	-0.156850	1.000000	0.753432	0.237567	0.005778	-0.042055	-0.016233	0.008842	-0.026126	
<b>Bathrooms</b>	0.773638	-0.248055	-0.194897	0.753432	1.000000	0.303258	0.024483	-0.052069	-0.018758	-0.010331	0.016043	
<b>Balcony</b>	0.361521	-0.151663	0.059638	0.237567	0.303258	1.000000	-0.013977	-0.035568	0.011083	0.001216	0.148764	
<b>Status</b>	0.024622	-0.005987	0.004647	0.005778	0.024483	-0.013977	1.000000	-0.127069	-0.002824	-0.006372	0.008236	
<b>neworold</b>	-0.067268	0.057831	0.097625	-0.042055	-0.052069	-0.035568	-0.127069	1.000000	0.016201	-0.004731	0.013361	
<b>parking</b>	-0.013302	-0.036681	0.003697	-0.016233	-0.018758	0.011083	-0.002824	0.016201	1.000000	-0.007023	0.025280	
<b>Furnished_status</b>	-0.014781	-0.044814	-0.008245	0.008842	-0.010331	0.001216	-0.006372	-0.004731	-0.007023	1.000000	0.043218	
<b>Lift</b>	0.043350	-0.023581	0.077498	-0.026126	0.016043	0.148764	0.008236	0.013361	0.025280	0.043218	1.000000	
<b>type_of_building</b>	0.017414	-0.008844	-0.005782	0.010986	0.010026	-0.002599	-0.005648	-0.006983	0.022445	-0.007999	0.018421	
<b>Price_sqft</b>	0.322863	-0.200779	-0.543093	0.313647	0.330614	0.028140	0.008838	-0.074156	-0.018500	-0.052717	-0.069917	

In [29]: `final_df.head()`

Out[29]:

	area	latitude	longitude	Bedrooms	Bathrooms	Balcony	Status	neworold	parking	Furnished_status	Lift	type_of_building	Price_sqft	
<b>0</b>	1350.0	28.608850	77.460560	3.0	3.0	2.0	1	0	1.0		1	2.0	0 4148.148148	
<b>1</b>	1490.0	28.374236	76.952416	3.0	3.0	2.0	0	0	1.0		1	2.0	0 5906.040268	
<b>2</b>	2385.0	28.645769	77.385110	4.0	5.0	2.0	0	0	1.0		2	2.0	0 6918.238994	
<b>3</b>	1050.0	28.566914	77.436434	2.0	2.0	3.0	0	0	1.0		2	2.0	0 3628.571429	
<b>4</b>	1350.0	28.520732	77.356491		2.0	2.0	3.0	0	1	1.0		1	3.0	0 4592.592593

In [30]: `final_df.drop(["Status"],axis=1,inplace=True)`

In [31]: `final_df.shape`

Out[31]: `(7360, 12)`

In [32]: `x=final_df.drop("Price_sqft",axis=1)  
y=final_df["Price_sqft"]  
x=np.array(x)  
y=np.array(y)`

In [33]: `x_col=x.shape[1]  
x_col`

Out[33]: `11`

Splitting the whole dataset into train and test datasets

In [34]: `from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=42)`

Scaling the features

In [35]: `from sklearn.preprocessing import StandardScaler  
sc=StandardScaler()  
x_train=sc.fit_transform(x_train)  
x_test=sc.transform(x_test)`

In [36]: `import tensorflow as tf  
from tensorflow import keras  
from keras import Sequential`

```
from keras.layers import Dense, BatchNormalization, Dropout
import keras_tuner as kt
```

WARNING:tensorflow:From C:\Users\mdsk5\anaconda3\lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse\_softmax\_cross\_entropy is deprecated. Please use tf.compat.v1.losses.sparse\_softmax\_cross\_entropy instead.

Hyperparameter Tuning using RandomSearchCV

```
In [37]: def build_model(hp):
    model=Sequential()
    model.add(keras.Input(shape=(x_col,)))
    for i in range(hp.Int("layers",min_value=1,max_value=10)):
        model.add(
            Dense(
                hp.Int("unit"+str(i),min_value=5,max_value=300,step=60),
                activation=hp.Choice("activation"+str(i),values=["relu","tanh"])
            )
        )
    model.add(Dense(1,activation="linear"))
    model.compile(optimizer=tf.keras.optimizers.Adam(
        hp.Float('learning_rate', min_value=1e-4, max_value=1e-2, sampling='LOG')),
        loss="mse",metrics=["mae"])

    return model
```

```
In [38]: tuner=kt.RandomSearch(build_model,
                           objective="mae",
                           max_trials=20,
                           directory='my_dir',
                           project_name='hp')
```

Reloading Tuner from my\_dir\hp\tuner0.json

```
In [39]: tuner.search(x_train,y_train,epochs=20,validation_data=(x_test,y_test))
```

Trial 20 Complete [00h 00m 09s]  
mae: 2395.01318359375

Best mae So Far: 775.0404052734375  
Total elapsed time: 00h 03m 51s

```
In [40]: tuner.get_best_hyperparameters()[0].values
```

```
Out[40]: {'layers': 5,
          'unit0': 245,
          'activation0': 'relu',
          'learning_rate': 0.0002662432681570655,
          'unit1': 5,
          'activation1': 'tanh',
          'unit2': 5,
          'activation2': 'relu',
          'unit3': 245,
          'activation3': 'relu',
          'unit4': 245,
          'activation4': 'relu',
          'unit5': 185,
          'activation5': 'relu',
          'unit6': 125,
          'activation6': 'relu',
          'unit7': 185,
          'activation7': 'tanh',
          'unit8': 125,
          'activation8': 'tanh',
          'unit9': 5,
          'activation9': 'tanh'}
```

```
In [41]: model=tuner.get_best_models(num_models=1)[0]
model.summary()
```

```
WARNING:tensorflow:From C:\Users\mdsk5\anaconda3\Lib\site-packages\keras\src\saving\legacy\save.py:538: The name tf.train.NewCheckpointReader is deprecated. Please use tf.compat.v1.train.NewCheckpointReader instead.
```

Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
dense (Dense)	(None, 245)	2940
dense_1 (Dense)	(None, 5)	1230
dense_2 (Dense)	(None, 5)	30
dense_3 (Dense)	(None, 245)	1470
dense_4 (Dense)	(None, 245)	60270
dense_5 (Dense)	(None, 1)	246
<hr/>		
Total params: 66186 (258.54 KB)		
Trainable params: 66186 (258.54 KB)		
Non-trainable params: 0 (0.00 Byte)		

Using callbacks to prevent overfitting

```
In [42]: callback=keras.callbacks.EarlyStopping(  
    monitor="val_loss",  
    min_delta=0,  
    patience=5,  
    verbose=1,  
    mode="auto",  
    baseline=None,  
)
```

```
In [43]: history=model.fit(x_train,y_train,epochs=200,validation_data=(x_test,y_test),callbacks=callback)
```

Epoch 1/200  
184/184 [=====] - 1s 3ms/step - loss: 1050349.7500 - mae: 771.1406 - val\_loss: 1134551.2500 - val\_mae: 805.5239  
Epoch 2/200  
184/184 [=====] - 0s 2ms/step - loss: 1041084.1875 - mae: 766.4742 - val\_loss: 1124809.7500 - val\_mae: 806.0350  
Epoch 3/200  
184/184 [=====] - 0s 2ms/step - loss: 1030343.5625 - mae: 762.7241 - val\_loss: 1116075.7500 - val\_mae: 801.4805  
Epoch 4/200  
184/184 [=====] - 0s 2ms/step - loss: 1021626.1875 - mae: 759.6195 - val\_loss: 1137453.5000 - val\_mae: 801.6863  
Epoch 5/200  
184/184 [=====] - 0s 2ms/step - loss: 1015045.3750 - mae: 755.0263 - val\_loss: 1109939.8750 - val\_mae: 798.3193  
Epoch 6/200  
184/184 [=====] - 0s 2ms/step - loss: 1008032.6250 - mae: 752.6318 - val\_loss: 1099815.5000 - val\_mae: 793.0822  
Epoch 7/200  
184/184 [=====] - 0s 2ms/step - loss: 997602.4375 - mae: 749.6944 - val\_loss: 1103051.2500 - val\_mae: 792.4496  
Epoch 8/200  
184/184 [=====] - 0s 2ms/step - loss: 989874.0000 - mae: 747.2515 - val\_loss: 1091946.1250 - val\_mae: 783.2615  
Epoch 9/200  
184/184 [=====] - 0s 2ms/step - loss: 987049.1250 - mae: 745.0770 - val\_loss: 1085937.8750 - val\_mae: 788.6705  
Epoch 10/200  
184/184 [=====] - 1s 3ms/step - loss: 978434.8750 - mae: 742.0422 - val\_loss: 1075823.7500 - val\_mae: 780.5148  
Epoch 11/200  
184/184 [=====] - 0s 2ms/step - loss: 974809.5625 - mae: 740.3794 - val\_loss: 1077887.0000 - val\_mae: 785.3599  
Epoch 12/200  
184/184 [=====] - 0s 2ms/step - loss: 964707.1250 - mae: 734.7602 - val\_loss: 1072209.8750 - val\_mae: 783.3728  
Epoch 13/200  
184/184 [=====] - 0s 2ms/step - loss: 956347.9375 - mae: 731.6196 - val\_loss: 1059081.7500 - val\_mae: 775.7266  
Epoch 14/200  
184/184 [=====] - 0s 2ms/step - loss: 955677.8125 - mae: 733.4897 - val\_loss: 1069719.6250 - val\_mae: 769.1562  
Epoch 15/200  
184/184 [=====] - 0s 2ms/step - loss: 948833.3750 - mae: 728.4099 - val\_loss: 1052706.6250 - val\_mae: 774.6198  
Epoch 16/200  
184/184 [=====] - 0s 2ms/step - loss: 940407.5000 - mae: 726.6315 - val\_loss: 1061406.1250 - val\_mae: 767.3257  
Epoch 17/200  
184/184 [=====] - 0s 2ms/step - loss: 940823.5000 - mae: 725.0844 - val\_loss: 1054740.7500 - val\_mae: 777.8772  
Epoch 18/200  
184/184 [=====] - 0s 2ms/step - loss: 932829.5625 - mae: 724.0141 - val\_loss: 1052434.6250 - val\_mae: 777.4741  
Epoch 19/200  
184/184 [=====] - 0s 2ms/step - loss: 932989.1875 - mae: 723.5446 - val\_loss: 1040966.0625 - val\_mae: 763.0253  
Epoch 20/200  
184/184 [=====] - 0s 2ms/step - loss: 926398.0000 - mae: 718.0724 - val\_loss: 1037669.7500 - val\_mae: 764.5653  
Epoch 21/200  
184/184 [=====] - 0s 2ms/step - loss: 917791.5000 - mae: 717.8456 - val\_loss: 1036296.4375 - val\_mae: 759.4383  
Epoch 22/200  
184/184 [=====] - 0s 2ms/step - loss: 923847.5000 - mae: 719.1830 - val\_loss: 1048330.9375 - val\_mae: 765.0922  
Epoch 23/200  
184/184 [=====] - 0s 2ms/step - loss: 914069.8125 - mae: 714.0080 - val\_loss: 1031397.7500 - val\_mae: 763.8140  
Epoch 24/200  
184/184 [=====] - 0s 2ms/step - loss: 911961.1875 - mae: 713.8069 - val\_loss: 1031482.6875 - val\_mae: 756.2954  
Epoch 25/200  
184/184 [=====] - 0s 2ms/step - loss: 908059.6250 - mae: 712.2388 - val\_loss: 1026844.0625 - val\_mae: 761.4702  
Epoch 26/200  
184/184 [=====] - 0s 2ms/step - loss: 904484.1875 - mae: 710.7975 - val\_loss: 1041968.0000 - val\_mae: 764.5232

```

Epoch 27/200
184/184 [=====] - 0s 2ms/step - loss: 900236.6875 - mae: 709.2963 - val_loss: 1025207.7500 - val_ma
e: 761.3053
Epoch 28/200
184/184 [=====] - 0s 3ms/step - loss: 898880.9375 - mae: 709.0389 - val_loss: 1024822.1875 - val_ma
e: 762.2409
Epoch 29/200
184/184 [=====] - 0s 2ms/step - loss: 897594.4375 - mae: 708.5884 - val_loss: 1017260.8125 - val_ma
e: 758.2079
Epoch 30/200
184/184 [=====] - 0s 2ms/step - loss: 896890.9375 - mae: 706.1305 - val_loss: 1015632.0625 - val_ma
e: 752.6676
Epoch 31/200
184/184 [=====] - 0s 2ms/step - loss: 891738.1875 - mae: 706.9269 - val_loss: 1018317.5000 - val_ma
e: 753.8192
Epoch 32/200
184/184 [=====] - 0s 2ms/step - loss: 888681.9375 - mae: 706.4627 - val_loss: 1020068.0625 - val_ma
e: 759.7234
Epoch 33/200
184/184 [=====] - 0s 2ms/step - loss: 885374.9375 - mae: 704.3605 - val_loss: 1023940.6250 - val_ma
e: 761.7509
Epoch 34/200
184/184 [=====] - 0s 2ms/step - loss: 882950.8750 - mae: 702.7349 - val_loss: 1030001.6250 - val_ma
e: 757.2977
Epoch 35/200
184/184 [=====] - 0s 2ms/step - loss: 879903.6250 - mae: 702.5853 - val_loss: 1006441.6250 - val_ma
e: 748.2228
Epoch 36/200
184/184 [=====] - 0s 2ms/step - loss: 879767.5625 - mae: 701.7951 - val_loss: 1004853.6250 - val_ma
e: 747.8186
Epoch 37/200
184/184 [=====] - 0s 2ms/step - loss: 874734.0625 - mae: 700.2505 - val_loss: 1021284.2500 - val_ma
e: 753.0834
Epoch 38/200
184/184 [=====] - 0s 2ms/step - loss: 874325.3750 - mae: 699.5425 - val_loss: 1011568.4375 - val_ma
e: 756.8805
Epoch 39/200
184/184 [=====] - 0s 2ms/step - loss: 867895.0625 - mae: 697.3554 - val_loss: 1021232.5000 - val_ma
e: 752.7552
Epoch 40/200
184/184 [=====] - 0s 2ms/step - loss: 869066.8750 - mae: 698.2853 - val_loss: 1009981.7500 - val_ma
e: 745.3506
Epoch 41/200
184/184 [=====] - 0s 2ms/step - loss: 867708.4375 - mae: 696.8190 - val_loss: 1008312.6875 - val_ma
e: 753.9169
Epoch 41: early stopping

```

Model evaluation

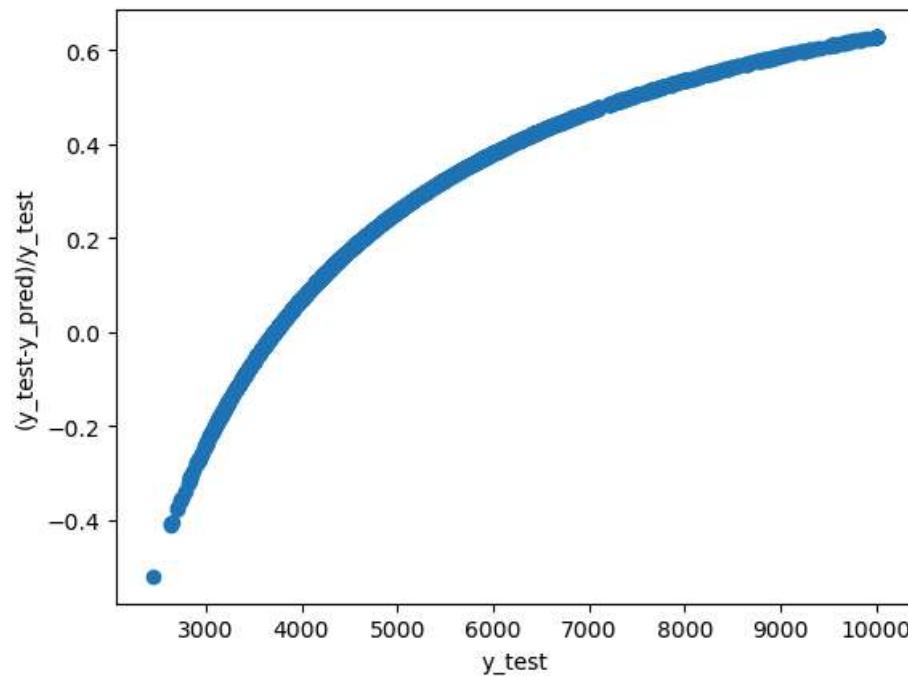
```
In [44]: from sklearn.metrics import mean_absolute_error,r2_score,accuracy_score

# Compute MAE
y_pred=model.predict(x_test)
mae = mean_absolute_error(y_test,y_pred )
r2=r2_score(y_test,y_pred)
print(f"Mean Absolute Error (MAE): {mae}")
print(f"r2 score: {r2}")

46/46 [=====] - 0s 1ms/step
Mean Absolute Error (MAE): 753.916823335624
r2 score: 0.6421548497867977
```

Plotting relative differences in output

```
In [45]: res=y_test-y_pred[0]
plt.scatter(y_test,res/y_test)
plt.xlabel("y_test")
plt.ylabel("(y_test-y_pred)/y_test")
plt.show()
```



Comparison the performance between train and validation datasets

```
In [46]: plt.plot(history.history["mae"],color="red",label="train")
plt.plot(history.history["val_mae"],color="blue",label="validation")
plt.title('mae')
plt.ylabel('mae')
plt.xlabel('Epoch')
plt.legend()
plt.show()
```

